

# DAY 4:DEVOPS

## 1. Install Docker Desktop

### For Windows:

1. Go to the official Docker website: <https://www.docker.com/products/docker-desktop/>
2. Click on "Download for Windows (WSL2)"
3. Run the installer after downloading
4. Follow the installation steps (enable WSL2 if prompted)
5. Restart your system if required
6. Open Docker Desktop from the Start Menu
7. Ensure Docker is running (you will see "Docker is running" status)

## 2. Create Folder Structure and Files:

Open Command Prompt, Terminal, or any shell:

```
mkdir microservices-project
```

```
cd microservices-project
```

```
mkdir order-service
```

```
mkdir user-service
```

```
cd order-service
```

```
touch Dockerfile app.py requirements.txt
```

```
cd ..
```

```
cd user-service
```

```
touch Dockerfile app.py requirements.txt
```

```
cd ..
```

## 3. Add Content to Files

```
order-service/app.py
```

```
from flask import Flask
```

```
app = Flask(__name__)
```

```
@app.route('/')
```

```
def home():
```

```
    return "Hello from Order Service"
```

```
if __name__ == '__main__':
```

```
    app.run(host='0.0.0.0', port=5000)
```

```
order-service/requirements.txt
```

```
flask
```

```
order-service/Dockerfile
```

```
FROM python:3.9-slim
```

```
WORKDIR /app
```

```
RUN pip install -r requirements.txt
```

```
EXPOSE 5000
```

```
CMD ["python", "app.py"]
```

```
user-service/app.py
```

```
from flask import Flask
```

```
app = Flask(__name__)
```

```
@app.route('/')
```

```
def home():
```

```
    return "Hello from User Service"
```

```
if __name__ == '__main__':
```

```
    app.run(host='0.0.0.0', port=5001)
```

```
user-service/Dockerfile
```

```
FROM python:3.9-slim
```

```
WORKDIR /app
```

```
RUN pip install -r requirements.txt
```

## EXPOSE 5001

CMD ["python", "app.py"]

### 4. Build Docker Images

```
docker run -d -p 5000:5000 --name order-service-container order-service-image
```

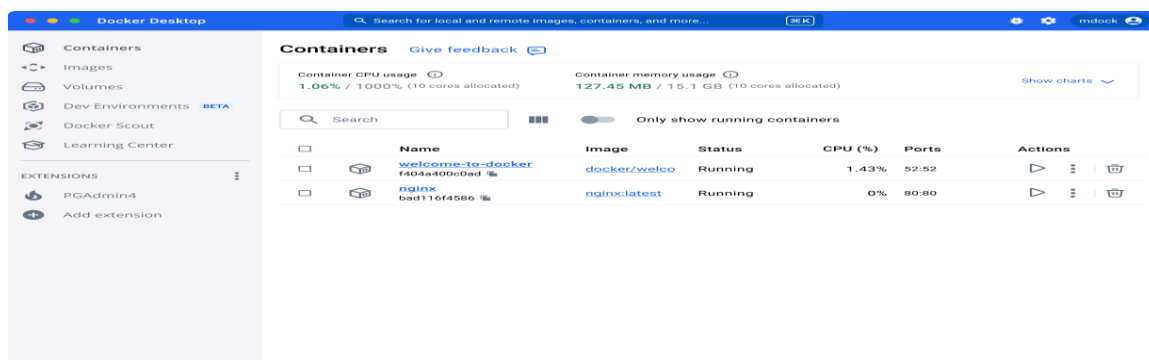
```
docker run -d -p 5001:5001 --name user-service-container user-service-image
```

### 6. Verify in Docker Desktop

1. Open Docker Desktop
2. Go to the "Containers" tab
3. You will see both order-service-container and user-service-container listed and running

```
Apr 24 microservices-project/order-service$ docker build -t
order-service-image ./order-service
=> muid image 1e90d833
=> => writing image from 309d8e3343d3e689f5eb9fce2033e5a6ea775b3f5a3
2ea02b1e22

Apr 24 microservices-project/user-service$ docker build -t
user-service-image ./user-service
=> build image 63fe35d9a79
=> => writing image from 65fe339946376889d3a6f61075955ab0fed9e92b127
441d3a66
```



## Creating an S3 Bucket in the AWS Console

1. Sign in to the AWS Management Console and open the S3 service.
2. Click “Create bucket.”
3. Enter a Bucket name (must be globally unique) and choose an AWS Region.
4. (Optional) Under Bucket settings for Block Public Access, leave defaults to block public access (recommended).

5. (Optional) Under Versioning, enable if you need object versioning.
6. (Optional) Under Encryption, choose AWS-managed key (SSE-S3) or your own KMS key.
7. (Optional) Add tags in the Tags section.
8. Review settings and click “Create bucket.”

## 2. Creating an S3 Bucket with Terraform

Prerequisites:

- Terraform installed (v1.0+).
- AWS credentials configured (e.g. via ~/.aws/credentials or environment variables).

In an empty project folder, create a file main.tf with:

```
terraform {  
  required_providers {  
    aws = {  
      source = "hashicorp/aws"  
      version = "~> 4.0"  
    }  
  }  
  required_version = ">= 1.0.0"  
}  
  
provider "aws" {  
  region = "us-east-1"      # change to your region  
}  
  
resource "aws_s3_bucket" "my_bucket" {  
  bucket = "my-unique-bucket-123" # change to a globally unique name  
  acl    = "private"  
  tags = {  
    Environment = "dev"  
    CreatedBy   = "Terraform"  
  }  
}
```

}

}

### Initialize Terraform:

terraform init

### Review the plan:

terraform plan

### Apply the plan:

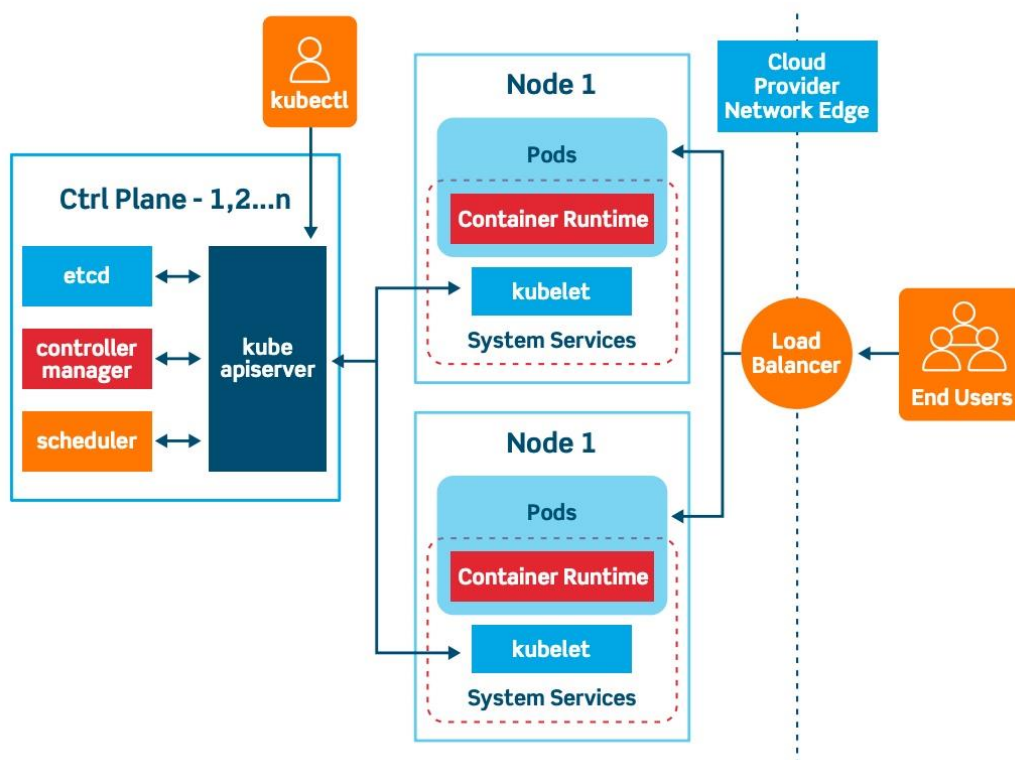
terraform apply

### (When done) To destroy:

terraform destroy

## Introduction to Kubernetes

Kubernetes is an orchestration platform—a system or tool that automates deployment, scaling, management, and operation of containerized applications. Note: Kubernetes uses Docker (or another container runtime) internally to run containers. With Kubernetes (often abbreviated “K8s”), you manage your Docker containers declaratively.



## Docker Swarm vs. Kubernetes

- Auto-scaling
  - Docker Swarm: scaling must be triggered manually (docker service scale ...).
  - Kubernetes: supports Horizontal Pod Autoscaler (HPA) to scale pods based on CPU, memory, or custom metrics.
- Production readiness
  - Docker Swarm: simpler, but fewer enterprise-grade features.
  - Kubernetes: richer feature set (self-healing, rolling updates, auto-scaling, namespaces) and is the industry standard for production.
- Recommendation

For production deployments, Kubernetes is highly recommended. It is effectively the successor to Docker Swarm for orchestrating containers at scale.

## Auto-Scaling

Auto-scaling automatically increases or decreases the number of running containers (pods) based on observed load (e.g., CPU usage, request rate). In Kubernetes, this is implemented via the Horizontal Pod Autoscaler.

## What Is a Kubernetes Cluster?

- Cluster: a group of servers (nodes) managed together.
  - Master Node(s): control plane components (API server, scheduler, controller manager) that accept user/developer instructions.
  - Worker Node(s): run the containerized applications as pods.
- Workflow:
  1. You (DevOps engineer or developer) submit a manifest (deployment, service, etc.) to the Kubernetes API on the master node.
  2. The master node schedules pods onto worker nodes.
  3. Worker nodes run your containers inside pods.
- High Availability: by distributing pods across multiple nodes, Kubernetes ensures your application stays available even if individual nodes fail.

Installing kubectl on Linux

# Download the binary

```
curl -LO "https://dl.k8s.io/release/$(curl -L -s https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubectl"
```

# Make it executable

```
chmod +x kubectl
```

# Move to PATH

```
sudo mv kubectl /usr/local/bin/
```

```
# Verify
```

```
kubectl version --client
```

Installing AWS CLI

```
curl "https://awscli.amazonaws.com/awscli-exe-linux-x86_64.zip" -o "awscliv2.zip"
```

```
unzip awscliv2.zip
```

```
sudo ./aws/install
```

```
aws --version
```

Installing Git

```
sudo yum install git -y # on Amazon Linux / CentOS
```

```
git --version # e.g., git version 2.47.1
```

**Clone the Voting App Repository**

```
git clone https://github.com/N4si/K8s-voting-app.git
```

```
cd K8s-voting-app
```

```
ls -l
```

```
cd manifests
```

```
ls
```

```
cd ..
```

**aws configure** **Connect to an EKS Cluster** **aws configure**

```
aws configure
```

**Update kubeconfig for your EKS cluster**

```
aws eks update-kubeconfig --name anvitha-cluster --region ap-south-1
```

**Verify nodes**

```
kubectl get nodes
```

# Kubernetes Architecture

